

```

# Import dependencies
import numpy as np
import torch
import torchvision
from torch.utils.data.dataset import Dataset
from torchvision import datasets, transforms
from torch import nn, optim
import matplotlib.pyplot as plt

```

### Problem 4.3 Compare different activation functions (sigmoid, Tanh, and ReLU)

```

# ##### Part 1: Load data and create batch
#####
N_total = 600
N_train = 500
x = torch.unsqueeze(torch.linspace(0, 1, N_total), dim=1)
r = torch.randperm(N_total)
x = x[r, :]
y = 0.2 + 0.4 * torch.pow(x, 2) + 0.3 * x * torch.sin(15 * x) + 0.05 *
    torch.cos(50 * x)

class CustomDataset(Dataset):
    def __init__(self, x, y):
        self.y = y
        self.x = x

    def __len__(self):
        return len(self.y)

    def __getitem__(self, idx):
        y1 = self.y[idx]
        x1 = self.x[idx]
        return (x1, y1)

# Change batch_size here to test different values
batch_size = 32 # Experiment with different batch sizes: 32, 64, 128
trainset = CustomDataset(x[0:N_train, :], y[0:N_train, :])
testset = CustomDataset(x[N_train:N_total, :], y[N_train:N_total, :])
train_loader = torch.utils.data.DataLoader(trainset,
    batch_size=batch_size)
test_loader = torch.utils.data.DataLoader(testset,
    batch_size=batch_size)

# ##### Part 2: Define Model and initialize
#####
def init_weights(m):
    if isinstance(m, nn.Linear):

```

```

        m.weight.data.uniform_(-1, 1)
        m.bias.data.uniform_(-1, 1)

# ##### Part 3: Define Loss and optimizer #####
criterion = torch.nn.MSELoss()

# ##### Part 4: Train and Test Functions #####
def train_NN(optimizer, model):
    model.train()
    for images, labels in train_loader:
        out = model(images)
        loss = criterion(out, labels)
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()
    return loss.item()

def test_NN(model, loader):
    model.eval()
    loss = 0
    with torch.no_grad():
        for images, labels in loader:
            out = model(images)
            loss += criterion(out, labels).item()
    loss /= len(loader)
    return loss

# ##### Compare Activation Functions #####
activation_functions = {
    "ReLU": nn.ReLU(),
    "Sigmoid": nn.Sigmoid(),
    "Tanh": nn.Tanh()
}

N_epoch = 400 # You can adjust the number of epochs

for activation_name, activation in activation_functions.items():
    # Define the model with the current activation function
    model = nn.Sequential(
        nn.Linear(1, 1024, bias=True),
        activation, # Use the current activation function
        nn.Linear(1024, 1, bias=True)
    )

    # Initialize weights
    model.apply(init_weights)

    # Initialize optimizer

```

```

optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# Initialize loss trackers
train_loss = np.zeros((N_epoch, 1))
test_loss = np.zeros((N_epoch, 1))

print(f"Testing {activation_name} activation function")

# Train the model
for epoch in range(N_epoch):
    train_loss[epoch, 0] = train_NN(optimizer, model)
    test_loss[epoch, 0] = test_NN(model, test_loader)
    if epoch % 100 == 0:
        print(f'Epoch: {epoch:03d}, Train Loss: {train_loss[epoch, 0]:.7f}, Test Loss: {test_loss[epoch, 0]:.7f}')

# Final prediction and plot
x_test = torch.unsqueeze(torch.linspace(0, 1, 1999), dim=1)
y_test = model(x_test)

# Plot the results
plt.figure()
plt.plot(x[0:N_total], y[0:N_total], 'bo', label='True data')
plt.plot(x_test, y_test.detach().numpy(), 'r', label=f'Model prediction ({activation_name})')
plt.legend()
plt.title(f"Prediction with {activation_name} activation")
plt.show()

# Plot train and test loss
plt.figure()
plt.plot(train_loss, label=f'Train Loss ({activation_name})')
plt.plot(test_loss, label=f'Test Loss ({activation_name})')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title(f'Train and Test Loss for {activation_name}')
plt.grid(True)
plt.show()

```

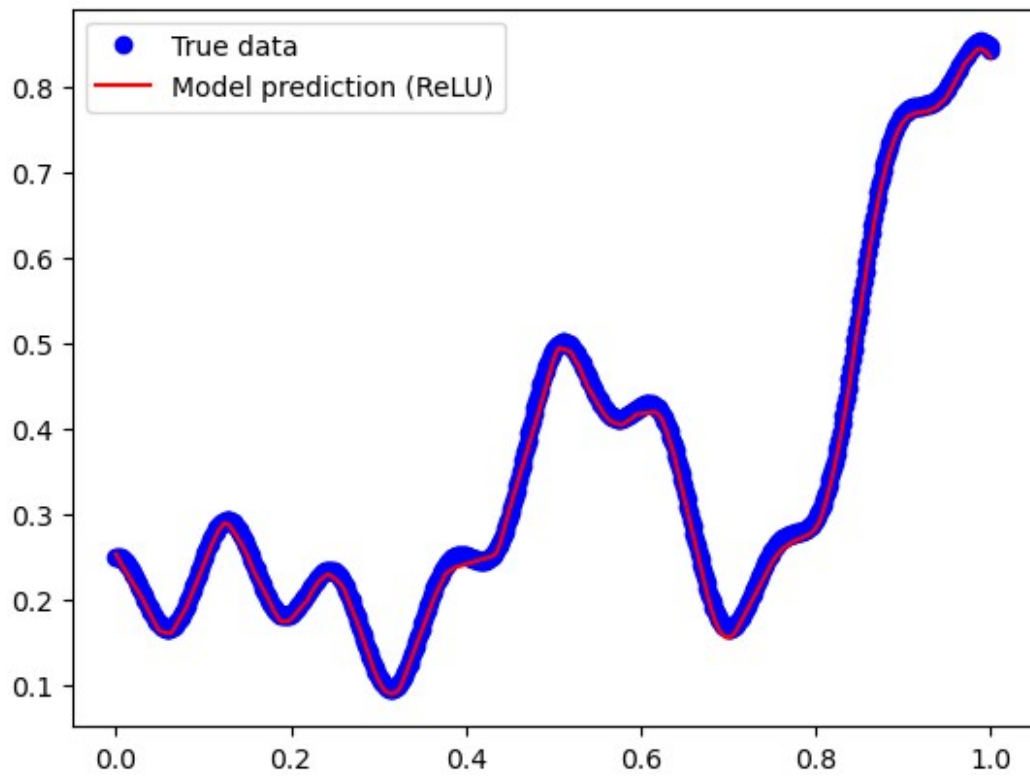
Testing ReLU activation function


```

Epoch: 000, Train Loss: 0.8893986, Test Loss: 1.0934367
Epoch: 100, Train Loss: 0.0005511, Test Loss: 0.0007331
Epoch: 200, Train Loss: 0.0003145, Test Loss: 0.0001574
Epoch: 300, Train Loss: 0.0000522, Test Loss: 0.0000628

```

Prediction with ReLU activation






Testing Sigmoid activation function


Epoch: 000, Train Loss: 0.6068853, Test Loss: 0.9664480

Epoch: 100, Train Loss: 0.0169263, Test Loss: 0.0201778

Epoch: 200, Train Loss: 0.0152558, Test Loss: 0.0202157

Epoch: 300, Train Loss: 0.0119938, Test Loss: 0.0173541





Testing Tanh activation function


Epoch: 000, Train Loss: 0.1610926, Test Loss: 0.1142495

Epoch: 100, Train Loss: 0.0116777, Test Loss: 0.0320398


Epoch: 200, Train Loss: 0.0148498, Test Loss: 0.0335571

Epoch: 300, Train Loss: 0.0182647, Test Loss: 0.0340703

Prediction with Tanh activation







Compare different activation functions (sigmoid, Tanh, and ReLU): ReLU